

Kubernetes @ NCSA

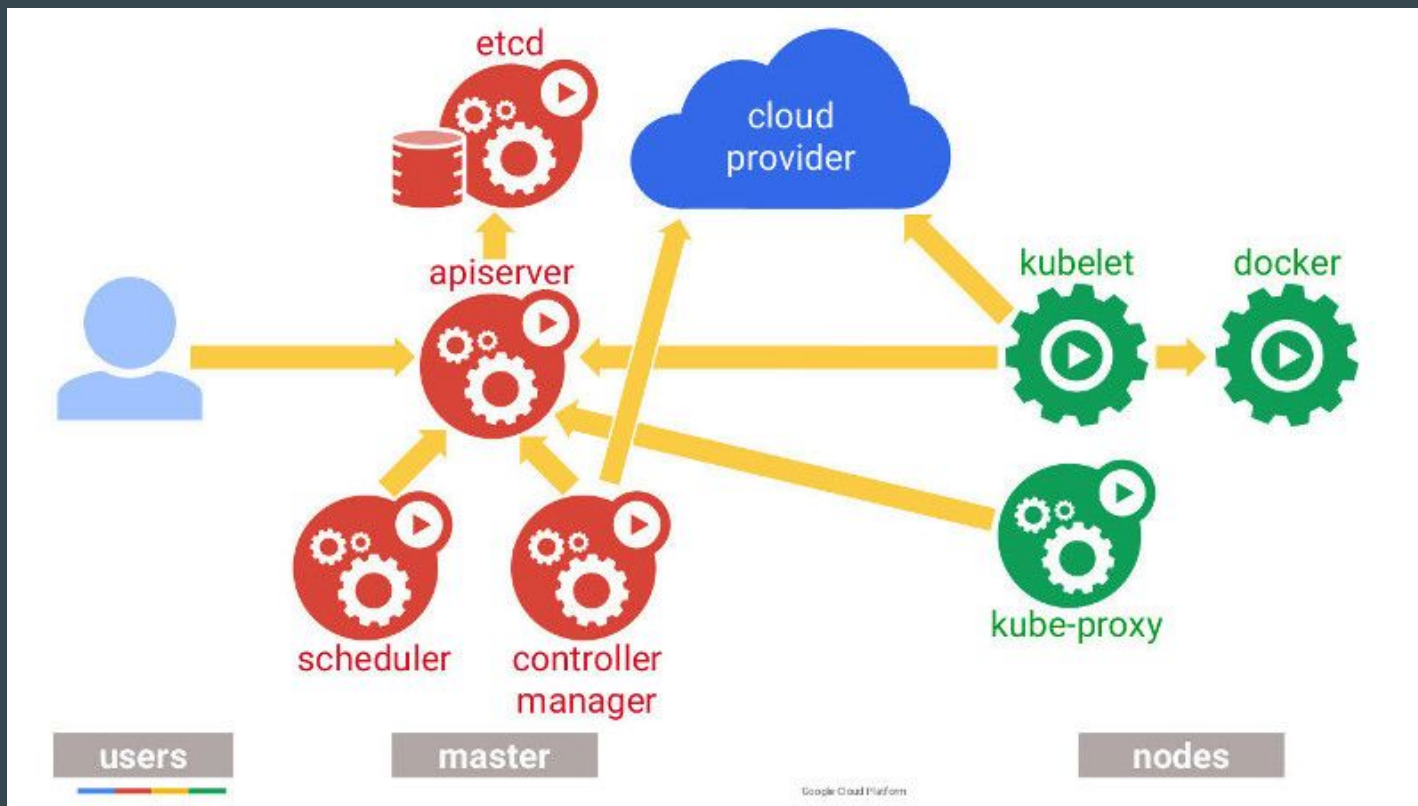


Matias Carrasco Kind • 09.14.2017

Outline

- Introduction
 - Basic concepts
 - Current efforts
 - @ DES
 - @ Nebula / Limitations
 - @ vSphere machines
 - Openshift
 - Storage
 - JupyterLab
 - Spark/Condor/Celery/Jobs
 - Security - RBAC
 - Namespaces
 - Users - Roles, ClusterRoles
 - ServiceAccounts ~ Roles/Cluster
 - Pod Policy
 - Network Policy
 - Resource Manager
 - Resource Quota
 - Selector Nodes
-
- Scalability
 - Monitor

Introduction

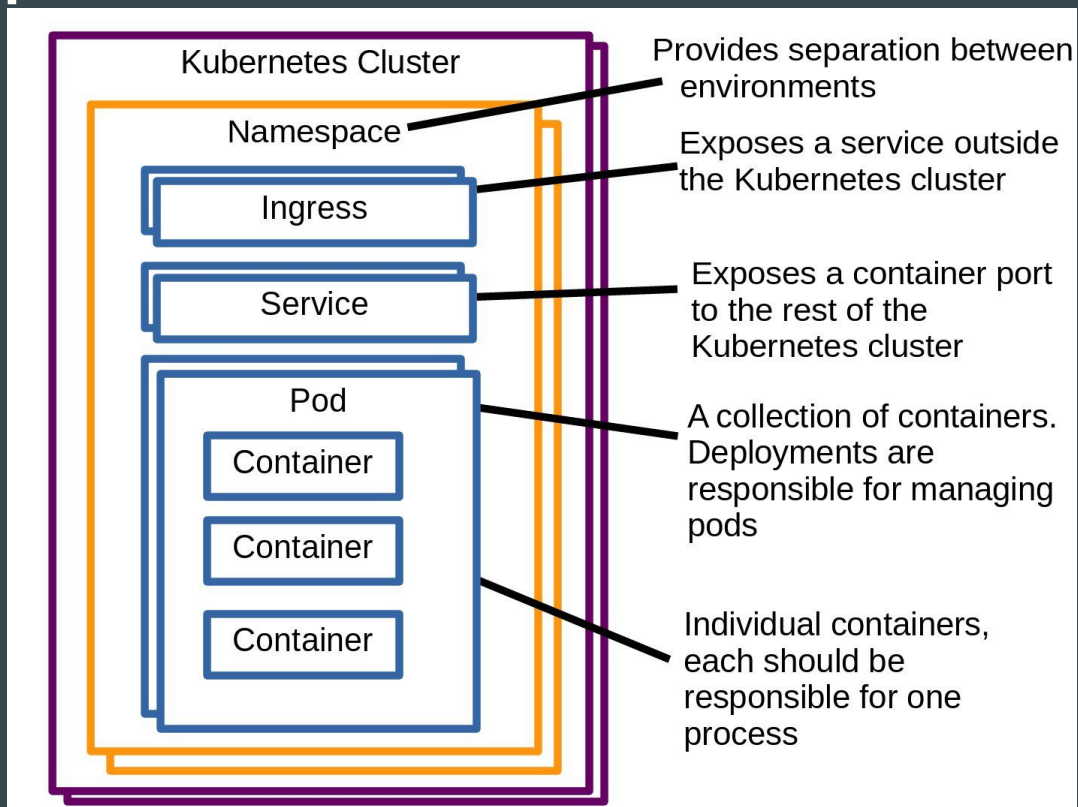


Introduction

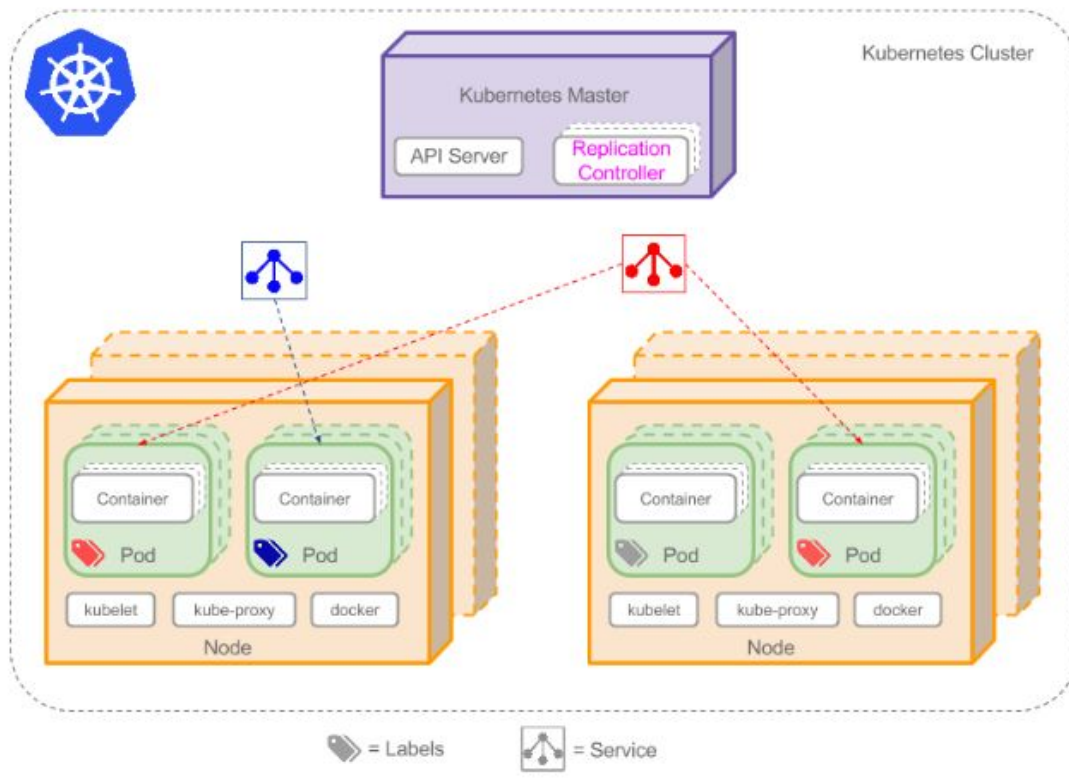
Key Concepts of Kubernetes

- **Pod** - A group of Containers
- **Labels** - Labels for identifying pods
- **Kubelet** - Container Agent
- **Proxy** - A load balancer for Pods
- **etcd** - A metadata service
- **cAdvisor** - Container Advisor provides resource usage/performance statistics
- **Replication Controller** - Manages replication of pods
- **Scheduler** - Schedules pods in worker nodes
- **API Server** - Kubernetes API server

Introduction



Introduction



Current Efforts @ NCSA

- DES Labs (deslabs.ncsa.illinois.edu), collection of services for DES, @Nebula
 - Frontpage
 - SQL Web client with front-end, monitor, Redis and job submission (pods = celery workers)
 - JupyterHub for internal DES
 - Cutout server
- DES Data Release interface, similar to DES Labs, using GPFS
- LSST K8s @ Nebula (lsstlabs.ncsa.illinois.edu) under development
 - 15 nodes (8 cores, 16 GB RAM and 160GB local disk per node)
 - 3 Cinder volumes attached to master (10TB, 1TB, 1TB) served using NFS
 - Nginx Ingress controller, RBAC,
- 2nd k8s lsst cluster @Nebula for testing
 - Spark, dns testing, namespaces cross services, celery workers, Jobs
- Kubernetes cluster in Vsphere machines, openshift → under R&D
- Other efforts @NCSA using kubernetes from Openstack/Nebula from other groups, monthly kubernetes meeting

Current Efforts @ NCSA : Limitations (to be addressed...)

- Openstack/Nebula \neq AWS, GKE
- No LoadBalancer type for services (dynamic ip provider), only using Ingress Rules, or especial ports
 - Not to be confused with Load Balancing for the end points of the services
- Cinder Volumes can only be attached to individual nodes, need to be served via NFS for Persistent volumes
- No dynamic volume provisioning
- Nodes need to be created manually, for scalability
- No read-write GPFS, read-only access using 2 NFS bridges (unless mounted to network)
- And other small technical issues...

Security

Namespaces/Labels

- Namespaces; can partition cluster in resources, users, etc. Different namespaces for different environments (prod vs devel)
- Labels: Used to select resources within the cluster or namespace, to select pods, nodes, deployments,

Users/Groups/Service Accounts

- User and Groups refer to humans running and using resources. Permissions/Roles are applied at these levels . Cluster scoped
- Service Accounts are for processes, permissions/Roles can be applied to allow a running pod to schedule another pod, etc... Namespace scoped

Security

Roles/Cluster Roles

- Roles are namespace scoped
 - CRUD resources namespace
 - Pods, deployments, PVC, service
 - Roles are bind to users/groups/sa
- Cluster Roles are cluster scoped
 - CRUD resources at cluster level
 - Nodes, namespaces, secrets, policy

Policies

- Pod Policy: at Cluster level to control how the pods/containers are run
 - Disabled running as root or a particular group, allow certain volumes to be mounted, limit access to port in host machine
- Network policy: how groups of pods are allowed to communicate with each other and other network endpoints. Namespaced scope, traffic control, use labels

Resource Manager

Resource Quotas

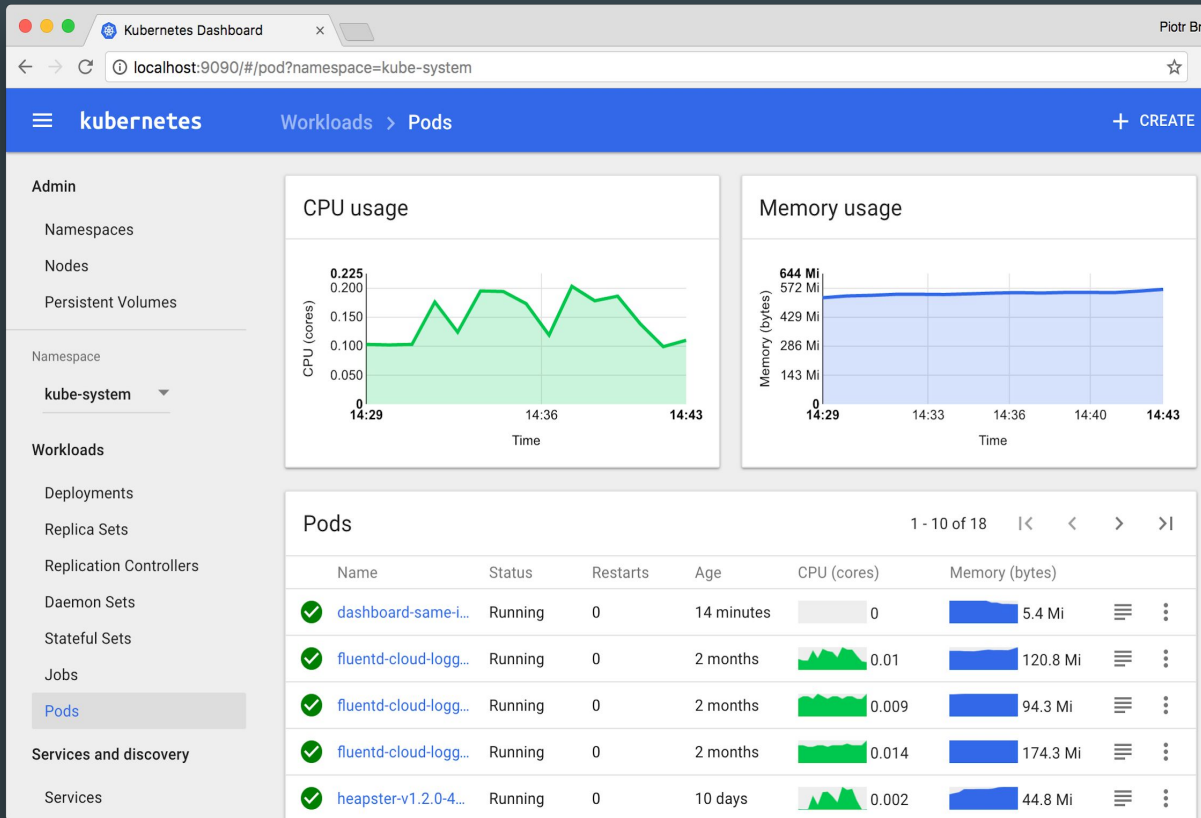
- Namespace scoped
- Limit cpu, memory, storage, etc
- Limit count of pods, deployments, claims, services, pods, etc.
- Doesn't work on nodes
- Can be updated dynamically

Node Selectors/Admission Control

- By labeling nodes, nodes can be tainted, reserved or specifically selected for scheduling
- Can enforce a set of dedicated nodes for a namespace using Admission Control
- Need to restart api server (not cluster)

Monitor

- Daemons sets can monitor nodes health, volumes
- Etcd @ master monitor use of resources, status of resources
- Controller manager monitors the status of the deployments and other resources
- Dashboard



Other features

Federation

- Allows to controls and manage multiple clusters

Scalability

- Up to 5,000 nodes and 150,000 pods
- Auto horizontal scaling can be enabled and scripted

Cluster

Roles (examples)

create/delete pods (R1)
Read pods (R2)
Create deployments (R3)
System:admin (R4)

Cluster Roles

List nodes (C1)
List namespaces (C2)
List secrets (C3)
System: admin (C4)

Network Policy

No traffic allow (N1)
Only pod-pod allow (N2)

Namespace 1 (sqre)

Resources

- RQ: cpu, memory, volume
- Limit selector to labeled nodes

Network Policy (N1)

Users:

🏠 R1,R2,R3,C1,C2



Service Account: SA_1 (R1)

Namespace 2 (devel)

Resources

- RQ: cpu, memory, volume
- Limit selector to labeled nodes

Network Policy (N1,N2)

Users:

💖 R1,R2,C1

Service Account: SA_1 (R1),
SA_2(C1,R2)

kube-system (admin)

😊 R4,C4

Users

Admin 😊

UserA 🏠

UserB 💖

Services Accounts

SA_1

SA_B

Pod Security Policy

-No root containers
-Limit volume
access

Resources

