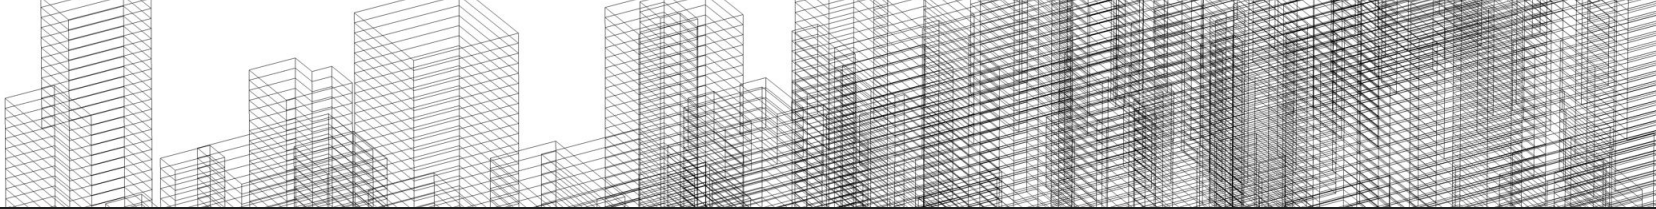


Batch and Online Anomaly Detection for Scientific Applications in a Kubernetes Environment

Matias Carrasco Kind
Senior Research Scientist, NCSA

Sahand Hariri
PhD Student, MechSE UIUC

ScienceCloud 18': 9th Workshop on Scientific Cloud Computing
June 11th, Tempe, AZ



Overview

Goal: Build a resilient scalable anomaly detection service.


Motivation: Astronomical data (both literal and figurative)

Algorithm: Extended Isolation Forest

Infrastructure: Kubernetes cluster


Mapreduce package: Spark

Part of the Motivation



Astronomy is just one example where data exploration needs to be automated.

Large catalogs, Large number of images, many unexpected objects/problems → Anomaly detection



- In operations 2020
- Every night for 10 years
- 18 billions objects (first year), ~40 billions by the end of survey
- ~1500 images per night
- Stream and static data
- Target to capture new physics (moving and variable objects)




DARK ENERGY SURVEY

- More than 500 nights of observation over 5 years
- 500 millions cataloged galaxies and 100 millions stars
- Many open problems: Systematics, new objects, new physics, etc.
- Almost completed



Anomaly Detection with Isolation Forest

- Few and different to be isolated quicker





Anomaly Detection with Isolation Forest

- Few and different to be isolated quicker
- For each tree:
 - Get a sample of the data





Anomaly Detection with Isolation Forest

- Few and different to be isolated quicker
- For each tree:
 - Get a sample of the data
 - Randomly select a dimension
 - Randomly pick a value in that dimension





Anomaly Detection with Isolation Forest

- Few and different to be isolated quicker
- For each tree:
 - Get a sample of the data
 - Randomly select a dimension
 - Randomly pick a value in that dimension
 - Draw a straight line through the data at that value and split data





Anomaly Detection with Isolation Forest

- Few and different to be isolated quicker
- For each tree:
 - Get a sample of the data
 - Randomly select a dimension
 - Randomly pick a value in that dimension
 - Draw a straight line through the data at that value and split data
 - Repeat until tree is complete






Anomaly Detection with Isolation Forest

- Few and different to be isolated quicker
- For each tree:
 - Get a sample of the data
 - Randomly select a dimension
 - Randomly pick a value in that dimension
 - Draw a straight line through the data at that value and split data
 - Repeat until tree is complete
- Generate multiple trees → forest





Anomaly Detection with Isolation Forest

- Few and different to be isolated quicker
- For each tree:
 - Get a sample of the data
 - Randomly select a dimension
 - Randomly pick a value in that dimension
 - Draw a straight line through the data at that value and split data
 - Repeat until tree is complete
- Generate multiple trees → forest
- Anomalies will be isolated in only a few steps



Anomaly Detection with Isolation Forest

- Few and different to be isolated quicker
- For each tree:
 - Get a sample of the data
 - Randomly select a dimension
 - Randomly pick a value in that dimension
 - Draw a straight line through the data at that value and split data
 - Repeat until tree is complete
- Generate multiple trees → forest
- Anomalies will be isolated in only a few steps
- Nominal points in more






Anomaly Detection with Isolation Forest

- Few and different to be isolated quicker
- For each tree:
 - Get a sample of the data
 - Randomly select a dimension
 - Randomly pick a value in that dimension
 - Draw a straight line through the data at that value and split data
 - Repeat until tree is complete
- Generate multiple trees → forest
- Anomalies will be isolated in only a few steps
- Nominal points in more


- To score points:

- Run point down tree, record path
- Repeat for each tree, aggregate scores
- Score distribution




Anomaly Detection with Isolation Forest



Single Tree scores for
anomaly and **nominal** points



Forest plotted radially.
Scores for **anomaly** and
nominal shown as lines





Anomaly Detection with Isolation Forest



Anomaly Detection with Extended Isolation Forest



Isolation Forest:

- ✓ Model free
- ✓ Computationally efficient
- ✓ Readily applicable to parallelization
- ✓ Readily application to high dimensional data
- ✗ Inconsistent scoring seen in score maps




Extended Isolation Forest:

- ✓ Model free
- ✓ Computationally efficient
- ✓ Readily applicable to parallelization
- ✓ Readily application to high dimensional data
- ✓ Consistent scoring



Anomaly Detection with Extended Isolation Forest



Algorithm 2 $iTree(X, e, l)$

Require: X - input data, e - current tree height, l - height limit


Ensure: an $iTree$

```

1: if  $e \geq l$  or  $|X| \leq 1$  then
2:   return  $exNode\{Size \leftarrow |X|\}$ 
3: else
4:   randomly select a normal vector  $n \in \mathbb{R}^{|X|}$  by drawing each coordinate of  $\vec{n}$  from a uniform
   distribution.
5:   randomly select an intercept point  $p \in \mathbb{R}^{|X|}$  in the range of  $X$ 
6:   set coordinates of  $n$  to zero according to extension level
7:    $X_l \leftarrow filter(X, (X - p) \cdot n \leq 0)$ 
8:    $X_r \leftarrow filter(X, (X - p) \cdot n > 0)$ 
9:   return  $inNode\{ Left \leftarrow iTree(X_l, e + 1, l),$ 
            $Right \leftarrow iTree(X_r, e + 1, l),$ 
            $Normal \leftarrow n,$ 
            $Intercept \leftarrow p\}$ 


```

10: **end if**



Technology Stack For Anomaly Service

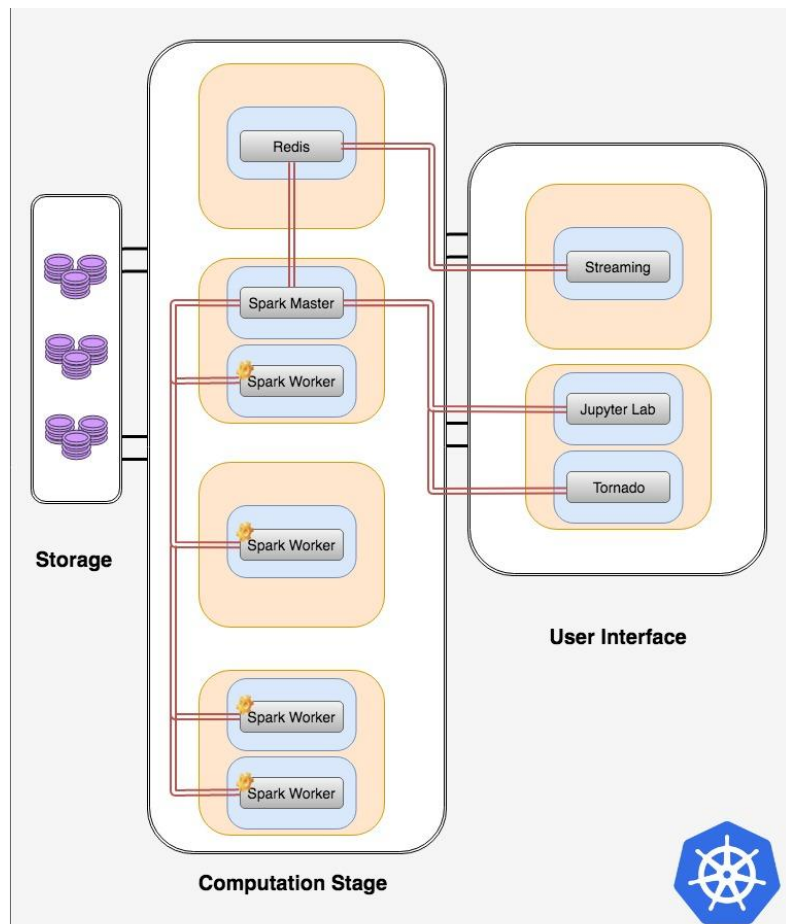
- Use Extended Isolation Forest as core algorithm
- Use Spark to parallelize trees and scoring
- Use Redis as a broker communicator
- To easily deploy in any environment, use Docker
- For orchestration of Docker containers, use Kubernetes
- Kubernetes cluster built on top of OpenStack, but it can be deployed also in AWS, GKE, etc.



Framework Architecture

There are three main components:

1. Storage
2. Computation Stage
3. User Interface / Streaming



Framework Architecture

Storage:

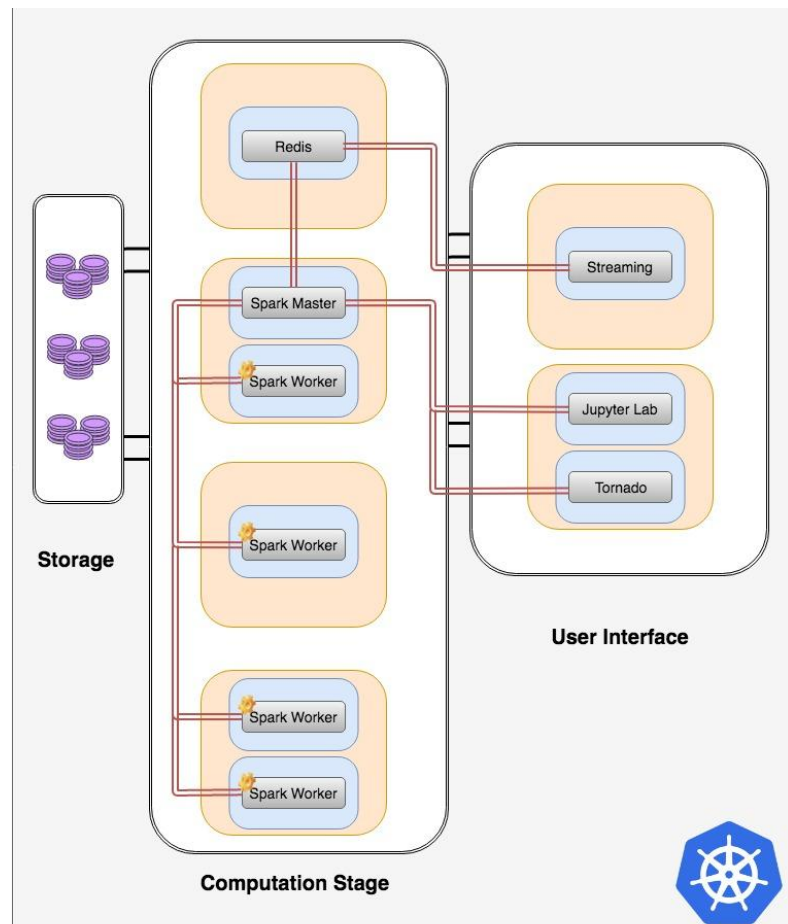
- NFS (Kubernetes PV/PVC)
- Redis
- RDD for Trees and Spark

User Interface:

- Jupyter notebooks
- Interactive web app for submitting jobs
- Streaming service

Computation Stage:

- Spark Master and Workers
- Communicator with Spark Master
- Suscription




Deployment

- Kubernetes allows very easy deployment, orchestration, scalability, resilience, replication, workloads and more
- From 0 to anomaly service → in minutes and config files
- Scale up/down (spark cluster and front-end) → Auto-scaling as an option
- Prototype support multiple users/projects, batch and streaming process
- Fault tolerant, disaster recovery




Spark Configuration Example

- Case 1: 800 trees, single core, serial mode
- Case 2: 100 trees on each core, aggregation and MapReduce. Each core access same data
- Case 3: Sample data on each core for 100 trees each, aggregation and MapReduce
- Case 4: Sample data on each core, 800 trees each
- Case 5: Sample data on each core, 100 trees on each, no aggregation




(a) Average difference computed as shown in equation 1



(b) Total average time taken to run each case.

Examples



Jupyter Notebooks

jupyter IFParallelExample Last Checkpoint: 4 minutes ago (autosaved) Logout

File Edit View Insert Cell Kernel Help Trusted Python 3

Create Spark Context

```
In [123]: from pyspark import SparkContext, SparkConf

In [124]: conf = SparkConf().setAppName("JupyterExamples").setMaster("spark://spark-master:7077")
          conf.set("spark.cores.max", 4)

Out[124]: <pyspark.conf.SparkConf at 0x7f7419428470>

In [134]: if sc:
          sc.stop()
          sc = SparkContext(conf=conf)
```

Imports

```
In [135]: import matplotlib.pyplot as plt
          import numpy as np
          from scipy.stats import multivariate_normal
          import random as rn
          import iso_forest as iso
          import seaborn as sb
          import time
          sb.set_style(style="whitegrid")
          sb.set_color_codes()
```

Helper Functions

```
In [136]: def getBlobData(N=2000):
          mean = [10, 1]
          cov = [[1, 0], [0, 1]] # diagonal covariance
          Nobjs = 4800
          x, y = np.random.multivariate_normal(mean, cov, Nobjs).T
          #Add manual outlier
          x[0]=3.3
          y[0]=3.3
          X=np.array([x,y]).T
          plt.figure(figsize=(7,7))
          plt.scatter(x,y,s=45,c=[0.5,0.5,0.5],alpha=0.3)
          plt.show()

          return (x,y,X)

In [137]: def getMultiBlobData(N=2000):
          mean1 = [10, 0]
          cov1 = [[1, 0], [0, 1]] # diagonal covariance
          mean2 = [0, 10]
          cov2 = [[1, 0], [0, 1]] # diagonal covariance
```

jupyter IFParallelExample Last Checkpoint: 5 minutes ago (autosaved) Logout

File Edit View Insert Cell Kernel Help Trusted Python 3

```
plt.plot(X[:,0],X[:,1], 'o', markersize=10, color=[0.5,0.5,0.5],alpha=0.3)
plt.axis("equal")

plt.show()

return (x,y,X)
```

```
In [138]: def getSinusoidData(N=4000):
          x = np.random.rand(N)*8*np.pi
          y = np.sin(x) + np.random.randn(N)/4.

          #Add manual outlier
          x[0]=3.3
          y[0]=3.3
          X=np.array([x,y]).T

          fig=plt.figure(figsize=(7,7))
          fig.add_subplot(111)
          plt.plot(X[:,0],X[:,1], 'o', markersize=10, color=[0.5,0.5,0.5], alpha=0.3)

          plt.show()

          return (x,y,X)
```

```
In [139]: def partition(l,n):
          return [l[i:i+n] for i in range(0,len(l),n)]
```

```
In [140]: def runIF(X):
          data = sc.parallelize(partition(X,int(len(X)/8)))
          Forest = data.map(lambda x: iso.iForest(x,ntrees=100, sample_size=256))
          S_t = Forest.map(lambda F: F.compute_paths(X))
          S = S_t.reduce(lambda a,b: a+b)/8
          return S
```

```
In [141]: def plotresults(x,y,scores):
          plt.rcParams['figure.figsize'] = (15, 5)
          plt.figure()
          plt.subplot(1,2,2)
          p=sb.distplot(scores, kde=True, color=[0.5,0.5,0.5])
          plt.xlabel('Anomaly Score',fontsize=20)
          plt.subplot(1,2,1)
          ss=np.argsort(scores)
          plt.scatter(x,y,s=45,c=[0.5,0.5,0.5],alpha=0.3)
          plt.scatter(x[ss[-10:]],y[ss[-10:]],s=55,c='r')
          plt.scatter(x[ss[:10]],y[ss[:10]],s=55,c='g')
          plt.show()
```

Examples

Jupyter Notebooks

Jupyter IParallelExample Last Checkpoint: 4 minutes ago (autosaved)


File Edit View Insert Cell Kernel Help Trusted Python 3

Code



Examples

Blob

```
In [148]: x,y,X = getBlobData()
```



```
In [149]: S = runIF(X)
plotresults(x,y,S)
```




Jupyter IParallelExample Last Checkpoint: 4 minutes ago (autosaved)

File Edit View Insert Cell Kernel Help Trusted Python 3



Code

MultiBlob

```
In [150]: x,y,X = getMultiBlobData()
```



```
In [151]: S = runIF(X)
plotresults(x,y,S)
```



Conclusions

- Open source anomaly detection software package for scientific application using fast and efficient isolation forest
- Fault tolerant, robust, scalable deployment
- Train and scoring using Spark
- Ready-to-deploy infrastructure on Kubernetes
- Production services for large datasets

Thank you!

Questions?

Matias Carrasco Kind -- NCSA

mcarras2@illinois.edu

github.com/mgkind

matias-ck.com

Extra Slides

Streaming

- 2 cases: Time evolving data, Time accumulative data
- Streaming isolation forest exists, not extended
- We can adapt and retrain trees as new data is presented
- Replace trees one by one until whole forest is replaced
- Work with window size to retrain trees